

## Assignment - PYSPARK

Affan Mohammed N Marikar  
281911

Questions:

- a) Create a new Spark Session with new SparkConfig

```
from pyspark import SparkConf, SparkContext
from pyspark.sql import SparkSession

config = SparkConf().setAppName("test_Session").setMaster("local[2]")
sc = SparkContext(conf = config)
spark = SparkSession.builder.appName('PySparkSession').getOrCreate()

sc
spark
```

- b) Create new instance of Spark SQL session and define new DataFrame using Flights\_Delay.csv dataset.

```
# Read the CSV file into a DataFrame
flights_delay_df = spark.read.csv("file:///home/hadoop/Downloads/Flights_Delay.csv",
header=True, inferSchema=True)
# Show the schema and first few rows of the DataFrame to verify
flights_delay_df.printSchema()
```

```
# Show the schema and first few rows of the DataFrame to verify
flights_delay_df.printSchema()
```

```
root
|-- ID: integer (nullable = true)
|-- YEAR: integer (nullable = true)
|-- MONTH: integer (nullable = true)
|-- DAY: integer (nullable = true)
|-- DAY_OF_WEEK: integer (nullable = true)
|-- AIRLINE: string (nullable = true)
|-- FLIGHT_NUMBER: integer (nullable = true)
|-- TAIL_NUMBER: string (nullable = true)
|-- ORIGIN_AIRPORT: string (nullable = true)
|-- DESTINATION_AIRPORT: string (nullable = true)
|-- SCHEDULED_DEPARTURE: integer (nullable = true)
|-- DEPARTURE_TIME: integer (nullable = true)
|-- DEPARTURE_DELAY: integer (nullable = true)
|-- TAXI_OUT: integer (nullable = true)
|-- WHEELS_OFF: integer (nullable = true)
|-- SCHEDULED_TIME: integer (nullable = true)
|-- ELAPSED_TIME: integer (nullable = true)
|-- AIR_TIME: integer (nullable = true)
|-- DISTANCE: integer (nullable = true)
```

flights\_delay\_df.show(5)

ID	YEAR	MONTH	DAY	DAY_OF_WEEK	AIRLINE	FLIGHT_NUMBER	TAIL_NUMBER	ORIGIN_AIRPORT	DESTINATION_AIRPORT	SCHEDULED_DEPARTURE	DEPARTURE_TIME	DEPARTURE_DELAY	TAXI_OUT	WHEELS_OFF	SCHEDULED_TIME	ELAPSED_TIME	AIR_TIME	DISTANCE	WHEELS_ON	TAXI_IN	SCHEDULED_ARRIVAL	ARRIVAL_TIME	ARRIVAL_DELAY	DIVERTED	CANCELLED	CANCELLATION_REASON	AIR_SYSTEM_DELAY	SECURITY_DELAY	AIRLINE_DELAY	LATE_AIRCRAFT_DELAY	WEATHER_DELAY	
935	2015	3	4	3	EV	5170	N842AS	CVG	XNA	129	1058	19	16	1010	115	129	108	562	14	1058	5	1030	1103	33	0	0	0	0	0	0	0	0
1240	2015	2	2	1	MQ	3584	N646MQ	DFW	SPS	46	1357	36	11	1327	50	46	30	113	0	1357	5	1330	1402	32	0	0	0	0	0	0	0	0
1335	2015	1	27	2	B6	716	N309JB	JAX	DCA	110	1652	90	16	1521	104	110	91	634	6	1652	3	1519	1655	96	0	0	0	0	0	0	0	0

c) Create table Spark HIVE table flights\_table

flights\_delay\_updated\_df.createOrReplaceTempView("flights\_table")

```
spark.sql("select * from flights_table").show()
```

ID	YEAR	MONTH	DAY	DAY_OF_WEEK	AIRLINE	FLIGHT_NUMBER	TAIL_NUMBER	ORIGIN_AIRPORT	DESTINATION_AIRPORT	SCHEDULED_DEPARTURE	DEPARTURE_TIME	DEPARTURE_DELAY	TAXI_OUT	WHEELS_OFF	SCHEDULED_TIME	ELAPSED_TIME	AIR_TIME	DISTANCE	WHEELS_ON	TAXI_IN	SCHEDULED_ARRIVAL	ARRIVAL_TIME	ARRIVAL_DELAY	DIVERTED	CANCELLED	CANCELLATION_REASON	AIR_SYSTEM_DELAY	SECURITY_DELAY	AIRLINE_DELAY	LATE_AIRCRAFT_DELAY	WEATHER_DELAY	
935	2015	3	4	3	EV	5170	N842AS	CVG	XNA	129	1058	19	16	1010	115	129	108	562	14	1058	5	1030	1103	33	0	0	0	0	0	0	0	0
0		19				0				null																						
1240	2015	2	2	1	MQ	3584	N646MQ	DFW	SPS	46	1357	36	11	1327	50	46	30	113	0	1357	5	1330	1402	32	0	0	0	0	0	0	0	0

d) Describe the table schema & show top 10 rows of Dataset

```
spark.sql("describe flights_table").show()
```

```
: spark.sql("describe flights_table").show()
```

col_name	data_type	comment
ID	int	null
YEAR	int	null
MONTH	int	null
DAY	int	null
DAY_OF_WEEK	int	null
AIRLINE	string	null
FLIGHT_NUMBER	int	null
TAIL_NUMBER	string	null
ORIGIN_AIRPORT	string	null
DESTINATION_AIRPORT	string	null
SCHEDULED_DEPARTURE	int	null
DEPARTURE_TIME	int	null
DEPARTURE_DELAY	int	null
TAXI_OUT	int	null
WHEELS_OFF	int	null
SCHEDULED_TIME	int	null
ELAPSED_TIME	int	null
AIR_TIME	int	null
DISTANCE	int	null
WHEELS_ON	int	null
TAXI_IN	int	null
SCHEDULED_ARRIVAL	int	null
ARRIVAL_TIME	int	null
ARRIVAL_DELAY	int	null
DIVERTED	int	null
CANCELLED	int	null
CANCELLATION_REASON	string	null
AIR_SYSTEM_DELAY	int	null
SECURITY_DELAY	int	null
AIRLINE_DELAY	int	null
LATE_AIRCRAFT_DELAY	int	null
WEATHER_DELAY	int	null

```
spark.sql("select * from flights_table limit 10").show()
```

```
spark.sql("select * from flights_table").show()
```

ID	YEAR	MONTH	DAY	DAY_OF_WEEK	AIRLINE	FLIGHT_NUMBER	TAIL_NUMBER	ORIGIN_AIRPORT	DESTINATION_AIRPORT	SCHEDULED_DEPARTURE	DEPARTURE_TIME	DEPARTURE_DELAY	TAXI_OUT	WHEELS_OFF	SCHEDULED_TIME	ELAPSED_TIME	AIR_TIME	DISTANCE	WHEELS_ON	TAXI_IN	SCHEDULED_ARRIVAL	ARRIVAL_TIME	ARRIVAL_DELAY	DIVERTED	CANCELLED	CANCELLATION_REASON	AIR_SYSTEM_DELAY	SECURITY_DELAY	AIRLINE_DELAY	LATE_AIRCRAFT_DELAY	WEATHER_DELAY
935	2015	3	4	3	EV	5170	N842AS	CVG	XNA	129	108	562	1058	14	No Reason	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
5	2015	3	4	3	EV	5170	N842AS	CVG	XNA	129	108	562	1058	14	No Reason	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
1240	2015	2	2	1	MQ	3584	N646MQ	DFW	SPS	46	30	113	1357	0	No Reason	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
5	2015	3	4	3	EV	5170	N842AS	CVG	XNA	129	108	562	1058	14	No Reason	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
1335	2015	1	27	2	B6	716	N309JB	JAX	DCA	110	91	634	1652	61	No Reason	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
3	2015	1	27	2	B6	716	N309JB	JAX	DCA	110	91	634	1652	61	No Reason	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0

- e) Apply Query performance optimization techniques like – creating Partitioning DataFrame by a specific column, parquet data, caching, predicate pushdown methods etc.

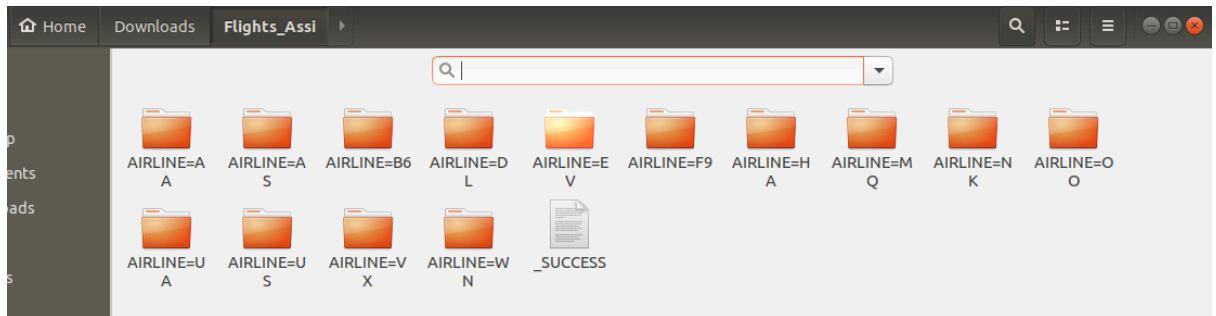
```
# Optimise by partitioning the DataFrame by a specific column (here; Airline)
```

```
partitioned_df = flights_delay_updated_df.repartition("AIRLINE")
```

```
# Save the DataFrame as Parquet files for better compression and faster queries
```

```
parquet_path = "file:///home/hadoop/Downloads/Flights_Assi"
```

```
partitioned_df.write.mode("overwrite").partitionBy("AIRLINE").parquet(parquet_path)
```



```
# Load the Parquet files into a DataFrame
```

```
flight_parq_df = spark.read.parquet(parquet_path)
```

```
# Create a Hive table using the optimised Parquet data
```

```
flight_parq_df.createOrReplaceTempView("flight_parq_hivetable")
```

```
spark.sql("CREATE DATABASE IF NOT EXISTS flights_db")
```

```
spark.sql("USE flights_db")
```

```
spark.sql("""
```

```
    CREATE TABLE IF NOT EXISTS flight_parq_table
```

```
    USING PARQUET
```

```
    LOCATION '{parquet_path}'
```

```
""").format(parquet_path=parquet_path))
```

```
spark.sql("CREATE DATABASE IF NOT EXISTS flights_db")
spark.sql("USE flights_db")
spark.sql("""
    CREATE TABLE IF NOT EXISTS flight_parq_table
    USING PARQUET
    LOCATION '{parquet_path}'
""").format(parquet_path=parquet_path))
```

```
DataFrame[]
```

```
# Cache the DataFrame for repeated queries
```

```
flight_parq_df.cache()
```

```
# Cache the DataFrame for repeated queries
flight_parq_df.cache()
```

```
DataFrame[ID: int, YEAR: int, MONTH: int, DAY: int, DAY_OF_WEEK: int, FLIGHT_NUMBER: int, TAIL_NUMBER: string, ORIGIN_AIRPORT: string, DESTINATION_AIRPORT: string, SCHEDULED_DEPARTURE: int, DEPARTURE_TIME: int, DEPARTURE_DELAY: int, TAXI_OUT: int, WHEELS_OFF: int, SCHEDULED_TIME: int, ELAPSED_TIME: int, AIR_TIME: int, DISTANCE: int, WHEELS_ON: int, TAXI_IN: int, SCHEDULED_ARRIVAL: int, ARRIVAL_TIME: int, ARRIVAL_DELAY: int, DIVERTED: int, CANCELLED: int, CANCELLATION_REASON: string, AIR_SYSTEM_DELAY: int, SECURITY_DELAY: int, AIRLINE_DELAY: int, LATE_AIRCRAFT_DELAY: int, WHEELS_ON_DELAY: int, AIRLINE: string]
```

```
# Enable predicate pushdown by filtering the DataFrame by filter data for ORIGIN_AIRPORT
```

```
filtered_df = flight_parq_df.filter("ORIGIN_AIRPORT = 'CVG'")
```

```
# Describe the optimised table schema
```

```
print("Optimised Table Schema:")
```

```
spark.sql("DESCRIBE flight_parq_table").show(truncate=False)
```

Optimized Table Schema:

col_name	data_type	comment
ID	int	null
YEAR	int	null
MONTH	int	null
DAY	int	null
DAY_OF_WEEK	int	null
FLIGHT_NUMBER	int	null
TAIL_NUMBER	string	null
ORIGIN_AIRPORT	string	null
DESTINATION_AIRPORT	string	null
SCHEDULED_DEPARTURE	int	null
DEPARTURE_TIME	int	null
DEPARTURE_DELAY	int	null
TAXI_OUT	int	null
WHEELS_OFF	int	null
SCHEDULED_TIME	int	null
ELAPSED_TIME	int	null
AIR_TIME	int	null
DISTANCE	int	null
WHEELS_ON	int	null

```
# Show the top 10 rows of the filtered dataset
```

```
print("Top 10 Rows of the Filtered Dataset:")
```

```
filtered_df.show(10, truncate=False)
```

Top 10 Rows of the Filtered Dataset:

```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|ID|YEAR|MONTH|DAY|DAY_OF_WEEK|FLIGHT_NUMBER|TAIL_NUMBER|ORIGIN_AIRPORT|DESTINATION_AIRPORT|SCHEDULED_DEPARTURE|D
EPARTURE_TIME|DEPARTURE_DELAY|TAXI_OUT|WHEELS_OFF|SCHEDULED_TIME|ELAPSED_TIME|AIR_TIME|DISTANCE|WHEELS_ON|TAXI_IN|SC
HEDULED_ARRIVAL|ARRIVAL_TIME|ARRIVAL_DELAY|DIVERTED|CANCELLED|CANCELLATION_REASON|AIR_SYSTEM_DELAY|SECURITY_DELAY|AI
RLINE_DELAY|LATE_AIRCRAFT_DELAY|WEATHER_DELAY|AIRLINE|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|767|2015|1|30|5|2525|N910DE|CVG|RSW|725|7|
20|-5|935|23|743|155|135|108|879|931|4|10
00|1143|0|0|null|DL|1579|2015|3|2|1|1143|N325US|CVG|ATL|600|6
01|1|20|621|91|98|68|373|729|10|73
1|1739|18|10|10|null|null|null|10|10|

```

f) Average arrival delay caused by airlines

```

avg_arrival_delay_airlines = spark.sql("""
    SELECT AIRLINE, round(AVG(CASE WHEN ARRIVAL_DELAY > 0 THEN ARRIVAL_DELAY
    ELSE NULL END),2) AS avg_arrival_delay
    FROM flights_table
    GROUP BY AIRLINE
    ORDER BY avg_arrival_delay DESC
""")
avg_arrival_delay_airlines.show()

```

```


+-----+-----+
|AIRLINE|avg_arrival_delay|
+-----+-----+
|F9|47.37|
|B6|42.78|
|MQ|42.57|
|NK|39.85|
|OO|37.54|
|DL|36.48|
|AA|36.29|
|EV|36.21|
|VX|35.14|
|UA|34.13|
|US|29.41|
|WN|27.64|
|AS|24.83|
|HA|16.05|
+-----+-----+

```

# Visualisation

```
avg_arrival_delay_pandas = avg_arrival_delay_airlines.toPandas()
```

```
avg_arrival_delay_pandas.plot(kind='bar', x='AIRLINE', y='avg_arrival_delay',
title='Average Arrival Delay by Airlines')
```



g) Days of months with respected to average of arrival delays


```
avg_arrival_delay_days = spark.sql("""
    SELECT DAY, round(AVG(CASE WHEN ARRIVAL_DELAY > 0 THEN ARRIVAL_DELAY
ELSE NULL END),2) AS avg_arrival_delay
    FROM flights_table
    GROUP BY DAY
    ORDER BY DAY
    """)
avg_arrival_delay_days.show(31)
```

DAY	avg_arrival_delay
1	43.55
2	37.13
3	41.83
4	42.05
5	42.09
6	39.52
7	31.25
8	34.96
9	31.23
10	24.73
11	33.65
12	36.25
13	28.53
14	27.39
15	26.97
16	33.15
17	37.84
18	27.11
19	24.6
20	29.92
21	34.86
22	31.69
23	26.55
24	33.48
25	33.48
26	37.02
27	36.13
28	33.38
29	25.16
30	32.82
31	31.97

# Visualisation

```
avg_arrival_delay_days_df = avg_arrival_delay_days.toPandas()
```

```
avg_arrival_delay_days_df.plot(kind='bar', x='DAY', y='avg_arrival_delay', title='Average  
Arrival Delay by Day of Month')
```



h) Arrange weekdays with respect to the average arrival delays caused

```
avg_arrival_delay_weekdays = spark.sql("""
    SELECT DAY_OF_WEEK, round(AVG(CASE WHEN ARRIVAL_DELAY > 0 THEN
    ARRIVAL_DELAY ELSE NULL END),2) AS avg_arrival_delay
    FROM flights_table
    GROUP BY DAY_OF_WEEK
    ORDER BY avg_arrival_delay DESC
""")
avg_arrival_delay_weekdays.show()
```


DAY_OF_WEEK	avg_arrival_delay
7	38.42
2	36.64
1	36.38
6	34.05
4	32.85
3	32.78
5	31.19

# Visualisation

```
avg_delay_weekdays_df = avg_arrival_delay_weekdays.toPandas()
```



```
avg_delay_weekdays_df.plot(kind='bar', x='DAY_OF_WEEK', y='avg_arrival_delay',
title='Average Arrival Delay by Day of Week')
```



i) Arrange Days of month as per cancellations done in Descending

```
cancellations_by_day = spark.sql("""
    SELECT DAY, COUNT(*) AS cancellations
    FROM flights_table
    WHERE CANCELLED = 1
    GROUP BY DAY
    ORDER BY cancellations DESC
    """)
```


```
cancellations_by_day.show(31)
```

DAY	cancellations
1	237
5	215
2	195
27	185
26	114
4	113
28	98
9	89
3	88
15	83
23	69
16	63
25	61
8	61
21	61
17	59
24	57
6	53
22	41
7	31
14	30
10	27
11	24
20	21
12	20
13	14
18	14
19	9
31	9
30	7
29	3

### # Visualisation

```
cancellations_by_day_df = cancellations_by_day.toPandas()
```

```
cancellations_by_day_df.plot(kind='bar', x='DAY', y='cancellations', title='Cancellations by Day of Month')
```



j) Find Top 10 busiest airports with respect to day of week

```
busiest_airports_by_day = spark.sql("""
SELECT Airport, DAY_OF_WEEK, SUM(flight_count) AS total_count FROM
( SELECT ORIGIN_AIRPORT AS Airport, DAY_OF_WEEK, COUNT(*) AS flight_count
  FROM flights_table GROUP BY DAY_OF_WEEK, ORIGIN_AIRPORT
  UNION ALL
  SELECT DESTINATION_AIRPORT AS Airport, DAY_OF_WEEK, COUNT(*) AS flight_count
  FROM flights_table GROUP BY DAY_OF_WEEK, DESTINATION_AIRPORT
)
AS combined GROUP BY Airport, DAY_OF_WEEK ORDER BY total_count DESC LIMIT 10
""")
busiest_airports_by_day.show()
```

Airport	DAY_OF_WEEK	total_count
ATL	5	1218
ATL	4	1113
ATL	1	1106
ATL	7	1021
ATL	3	985
ATL	2	960
ORD	5	925
DFW	5	881
ORD	1	844
ORD	4	832

#top 10 busiest airports:-

```
spark.sql("""select airportId, sum(total) as total_count from ( (select ORIGIN_AIRPORT as airportId, count(*) as total from flights_table group by ORIGIN_AIRPORT) UNION ALL (select DESTINATION_AIRPORT as airportId, count(*) as total from flights_table group by DESTINATION_AIRPORT) ) as combined group by airportId order by total_count DESC limit 10""").show()
```

airportId	total_count
ATL	7220
ORD	5655
DFW	5553
DEN	4088
LAX	4028
IAH	3228
PHX	3166
SFO	3031
LAS	2790
DTW	2374

k) Finding airlines that make the maximum number of cancellations


```
max_cancellations = spark.sql("""
    SELECT AIRLINE, COUNT(*) AS cancellation_count
    FROM flights_table
    WHERE CANCELLED = 1
    GROUP BY AIRLINE
    ORDER BY cancellation_count DESC
""")
max_cancellations.show()
```

AIRLINE	cancellation_count
MQ	414
WN	358
EV	312
AA	241
DL	177
US	169
OO	153
B6	145
UA	122
NK	21
VX	13
AS	12
F9	11
HA	3

# Visualisation

```
max_cancellations_df = max_cancellations.toPandas()
```

```
max_cancellations_df.plot(kind='bar', x='AIRLINE', y='cancellation_count', title='Max  
Cancellation VS Count')
```



- l) Find and order airlines in descending that make the most number of diversions

```
max_diversions_airlines = spark.sql("""
```

```
    SELECT AIRLINE, COUNT(*) AS diversion_count
```

```
    FROM flights_table
```

```
    WHERE DIVERTED = 1
```

```
    GROUP BY AIRLINE
```

```
    ORDER BY diversion_count DESC
```

```
""")
```


```
max_diversions_airlines.show()
```

```
+-----+
|AIRLINE|diversion_count|
+-----+
|      WN|              35|
|      OO|              25|
|      EV|              22|
|      DL|              18|
|      B6|              16|
|      AA|              12|
|      US|               9|
|      UA|               8|
|      MQ|               5|
|      HA|               1|
+-----+
```

# Visualisation

```
max_diversions_df = max_diversions_airlines.toPandas()
```

```
max_diversions_df.plot(kind='bar', x='AIRLINE', y='diversion_count', title='Max Diversion  
VS Count')
```



m) Finding days of month that see the most number of diversion

```
max_diversions_days = spark.sql("""
```

```
    SELECT DAY, COUNT(*) AS diversion_count
```

```
    FROM flights_table
```

```
    WHERE DIVERTED = 1
```

```
    GROUP BY DAY
```

```
    ORDER BY diversion_count DESC
```

```
""")
```


```
max_diversions_days.show()
```

DAY	diversion_count
2	15
1	13
4	12
5	11
9	9
14	8
6	7
23	6
7	6
8	5
3	5
11	5
30	5
18	5
28	4
15	28
13	16
12	20
11	21
9	12
8	31
7	26
6	17
27	27
13	13
19	19
10	10
15	15
25	25
29	29

# Visualisation

```
max_diversions_days_df = max_diversions_days.toPandas()
```

```
max_diversions_days_df.plot(kind='bar', y='DAY', x='diversion_count', title='Max Diversion  
VS Count')
```



n) Calculating mean and standard deviation of departure delay for all flights in minutes

```
filtered_flights_df = flights_delay_df.filter(col("DEPARTURE_DELAY") > 0)

departure_delay_stats = filtered_flights_df.select(
    round(mean("DEPARTURE_DELAY"),2).alias("mean_dep_del"),
    round(stddev("DEPARTURE_DELAY"),2).alias("std_dep_del")
)

departure_delay_stats.show()
```

```
+-----+-----+
|mean_dep_del|std_dep_del|
+-----+-----+
|      33.92|      54.23|
+-----+-----+
```

- o) Calculating mean and standard deviation of arrival delay for all flights in minutes

```
filtered_flights_df = flights_delay_df.filter(col("ARRIVAL_DELAY") > 0)

arrival_delay_stats = filtered_flights_df.select(
    round(mean("ARRIVAL_DELAY"),2).alias("mean_arv_del"),
    round(stddev("ARRIVAL_DELAY"),2).alias("std_arv_del")
)

arrival_delay_stats.show()
```

```
+-----+-----+
|mean_arv_del|std_arv_del|
+-----+-----+
|      34.61|      54.02|
+-----+-----+
```

- p) Finding all diverted Route from a source to destination Airport & which route is the most diverted

```
spark.sql("""select ORIGIN_AIRPORT, DESTINATION_AIRPORT, COUNT(*) as Route_Count
from flights_table where DIVERTED = 1 group by ORIGIN_AIRPORT, DESTINATION_AIRPORT
ORDER BY Route_Count DESC""").show()
```

```
+-----+-----+-----+
|ORIGIN_AIRPORT|DESTINATION_AIRPORT|Route_Count|
+-----+-----+-----+
|      HOU      |      DAL      |      2    |
|      PHL      |      SAN      |      2    |
|      STT      |      PHL      |      2    |
|      IAH      |      ASE      |      2    |
|      TPA      |      LGA      |      2    |
|      JFK      |      EGE      |      2    |
|      JFK      |      SEA      |      2    |
|      ORD      |      ASE      |      2    |
|      CLT      |      IAH      |      2    |
+-----+-----+-----+
```

- q) Finding AIRLINES with its total flight count, total number of flights arrival delayed by more than 30 Minutes, % of such flights delayed by more than 30 minutes when it is not Weekends with minimum count of flights from Airlines by more than 10. Also Exclude some of Airlines 'AK', 'HI', 'PR', 'VI' and arrange output in descending order by % of such count of flights.

```
spark.sql("""select AIRLINE, count(*) as Total_Flight_Count,
sum(case when ARRIVAL_DELAY > 30 then 1 else 0 end) as Delayed_Flight_Count,
round(100 * sum(case when ARRIVAL_DELAY > 30 and DAY_OF_WEEK not in (6,7)
then 1 else 0 end)/count(*),2) as PDelay
FROM flights_table
WHERE AIRLINE NOT IN ('AK', 'HI', 'PR', 'VI')
GROUP BY AIRLINE
HAVING COUNT(*) > 10
ORDER BY PDelay DESC
""").show()
```

AIRLINE	Total_Flight_Count	Delayed_Flight_Count	PDelay
F9	794	198	17.51
MQ	3502	775	17.16
B6	2548	485	14.13
NK	1048	186	13.26
EV	5916	874	11.24
OO	5708	859	11.09
UA	4701	653	10.57
AA	5250	700	9.22
VX	573	67	8.2
US	3925	452	7.9
DL	7989	746	7.41
WN	11738	1235	7.4
AS	1586	100	4.04
HA	722	38	3.19

- r) Finding AIRLINES with its total flight count with total number of flights departure delayed by less than 30 Minutes, % of such flights delayed by less than 30 minutes when it is Weekends with minimum count of flights from Airlines by more than 10. Also Exclude some of Airlines 'AK', 'HI', 'PR', 'VI' and arrange output in descending order by % of such count of flights.

```
spark.sql("""select AIRLINE, count(*) as Total_Flight_Count,
            sum(case when DEPARTURE_DELAY < 30 then 1 else 0 end) as
Delayed_Flight_Count,
            round(100 * sum(case when DEPARTURE_DELAY > 30 and DAY_OF_WEEK >5 then 1
else 0 end)/count(*),2) as PDelay
FROM flights_table
WHERE AIRLINE NOT IN ('AK', 'HI', 'PR', 'VI')
GROUP BY AIRLINE
HAVING COUNT(*) > 10
ORDER BY PDelay DESC
""").show()
```



AIRLINE	Total_Flight_Count	Delayed_Flight_Count	PDelay
F9	794	596	7.18
B6	2548	2092	4.75
NK	1048	859	4.48
MQ	3502	2848	4.11
AA	5250	4577	3.85
OO	5708	4883	3.59
UA	4701	4019	3.47
WN	11738	10299	3.37
VX	573	503	3.32
US	3925	3518	3.24
EV	5916	5117	3.06
AS	1586	1479	2.08
DL	7989	7187	2.04
HA	722	695	1.66

- s) When is the best time of day/day of week/time of a year to fly with minimum delays?

```
from pyspark.sql.functions import hour, minute, col, when, avg
```

```
# Convert SCHEDULED_DEPARTURE and SCHEDULED_ARRIVAL to hour of the day
flights_delay_updated_df =
flights_delay_updated_df.withColumn("SCHEDULED_DEPARTURE_HOUR",
(flights_delay_updated_df["SCHEDULED_DEPARTURE"] / 100).cast("int")) \
    .withColumn("SCHEDULED_ARRIVAL_HOUR",
(flights_delay_updated_df["SCHEDULED_ARRIVAL"] / 100).cast("int"))
```

```
# Calculate average departure and arrival delay by hour of the day
avg_delay_by_hour =
flights_delay_updated_df.groupBy("SCHEDULED_DEPARTURE_HOUR").agg(
    avg.when(col("DEPARTURE_DELAY") > 0,
col("DEPARTURE_DELAY")).alias("avg_departure_delay"),
    avg.when(col("ARRIVAL_DELAY") > 0,
col("ARRIVAL_DELAY")).alias("avg_arrival_delay")
).orderBy("SCHEDULED_DEPARTURE_HOUR")
```

```
# Show the results
avg_delay_by_hour.show()
```

SCHEDULED_DEPARTURE_HOUR	avg_departure_delay	avg_arrival_delay
0	27.8	23.6060606060606
1	35.588235294117645	25.68
2	13.75	11.0
3	34.0	51.4
4	20.333333333333332	15.0
5	30.97902097902098	30.914141414141415
6	36.09884467265725	31.341549295774648
7	31.937923250564335	29.16142735768904
8	34.25353283458021	33.96315028901734
9	33.173761946133794	32.386656557998485
10	32.0852314474651	32.733853797019165
11	31.459610027855152	33.02514367816092
12	34.24120234604106	34.89728096676737
13	31.47463768115942	32.56007509386733
14	33.801867911941294	35.57083042568039
15	34.75278810408922	37.62823061630219
16	32.515	34.525513585155736
17	35.29085140137494	38.11011235955056
18	36.0625	38.21972203838518
19	35.48476992871031	36.2907133243607

# Calculate average departure and arrival delay by day of the week

```
avg_delay_by_day_of_week = flights_delay_updated_df.groupBy("DAY_OF_WEEK").agg(
  avg(when(col("DEPARTURE_DELAY") > 0,
col("DEPARTURE_DELAY"))).alias("avg_departure_delay"),
  avg(when(col("ARRIVAL_DELAY") > 0,
col("ARRIVAL_DELAY"))).alias("avg_arrival_delay")
).orderBy("DAY_OF_WEEK")
```

# Show the results

```
avg_delay_by_day_of_week.show()
```

DAY_OF_WEEK	avg_departure_delay	avg_arrival_delay
1	36.272677744483794	36.381143162393165
2	36.54911489066296	36.64089121887287
3	32.74443227455276	32.78448867536033
4	32.00206307102859	32.852420991591764
5	30.34500683994528	31.191777041942604
6	32.45436893203883	34.050180650341225
7	37.09448574969021	38.41988267983945

# Calculate average departure and arrival delay by month

```
avg_delay_by_month = flights_delay_updated_df.groupBy("MONTH").agg(
  avg(when(col("DEPARTURE_DELAY") > 0,
col("DEPARTURE_DELAY"))).alias("avg_departure_delay"),
```

```

    avg(when(col("ARRIVAL_DELAY") > 0,
col("ARRIVAL_DELAY"))).alias("avg_arrival_delay")
).orderBy("MONTH")

```

```

# Show the results
avg_delay_by_month.show()

```

```

+-----+-----+-----+
|MONTH|avg_departure_delay|avg_arrival_delay|
+-----+-----+-----+
|    1|    32.58091594942089|32.91658105609205|
|    2|    34.15339616509587|35.21688172043011|
|    3|    36.873714285714286|37.72750790457028|
+-----+-----+-----+

```

- t) Which airlines are best airline to travel considering number of cancellations, arrival, departure delays and all reasons affecting performance of airline industry.

```

from pyspark.sql.functions import col, avg, sum

```

```

airline_performance = spark.sql("""

```

```

    SELECT

```

```

    AIRLINE,

```

```

    COUNT(*) AS total_flights,

```

```

    SUM(CANCELLED) AS total_cancellations,

```

```

    AVG(CASE WHEN DEPARTURE_DELAY > 0 THEN DEPARTURE_DELAY ELSE NULL END) AS avg_departure_delay,

```

```

    AVG(CASE WHEN ARRIVAL_DELAY > 0 THEN ARRIVAL_DELAY ELSE NULL END) AS avg_arrival_delay,

```

```

    AVG(CASE WHEN AIR_SYSTEM_DELAY > 0 THEN AIR_SYSTEM_DELAY ELSE NULL END) AS avg_air_system_delay,

```

```

    AVG(CASE WHEN SECURITY_DELAY > 0 THEN SECURITY_DELAY ELSE NULL END) AS avg_security_delay,

```

```

    AVG(CASE WHEN AIRLINE_DELAY > 0 THEN AIRLINE_DELAY ELSE NULL END) AS avg_airline_delay,

```

```

    AVG(CASE WHEN LATE_AIRCRAFT_DELAY > 0 THEN LATE_AIRCRAFT_DELAY ELSE NULL END) AS avg_late_aircraft_delay,

```

```

        AVG(CASE WHEN WEATHER_DELAY > 0 THEN WEATHER_DELAY ELSE NULL END) AS
avg_weather_delay

        FROM flights_table

        GROUP BY AIRLINE

        limit 5

""")

```

**# Top 5 Airline providers**

```

print("Top 5 best Airline providers are: ")

airline_performance.select('AIRLINE').show()

```

---

```

Top 5 best Airline providers are:
+-----+
|AIRLINE|
+-----+
|      UA|
|      NK|
|      AA|
|      EV|
|      B6|
+-----+

```